

Sex Ed Classification*

Insert Subtitle

Rosemarie Topp

July 24, 2026

Insert abstract

1 Introduction

Disclaimer: Due to the subject matter, this paper contains sensitive language that may be considered offensive. The state of sexuality education in schools varies wildly around the world. The topic of if, when, and how to introduce children to certain subjects is extremely divisive, and often raises religious, political, and moral issues. According to the World Health Organization (2026), comprehensive sexuality education is defined as “a curriculum-based, incremental and scientifically-accurate approach that teaches children and young people about the cognitive, emotional, physical and social aspects of sexuality in a culturally-relevant and **age-appropriate** way.” That being said, what can be considered age-appropriate for a given demographic is hard to define. Even if it can be universally agreed upon that a piece of content is inappropriate for children to consume, the exact reasons why may be hard to put into words. The exact criteria often comes down to what United States Supreme Court Justice Potter Stewart once said when asked to define obscene pornography: “I know it when I see it” (Lattman 2007).

In the age of social media, millions upon millions of posts go up each day, many of which contain sensitive themes. Websites such as Youtube or TikTok are frequented by curious children seeking educational information (Döring 2024). Back in the early 2000s, before censorship laws were put in place, the internet was more of a wild west. Users could post all manner of sensitive content, and as the internet became more accessible to younger audiences, it was not uncommon for those of a formative age to be exposed to shocking images and videos that were far from age-appropriate. Over the years, various countries have put legislation in place to help shelter younger users from sensitive content, such as Canada’s very recent Bill C-34, which mandates the implementation of safeguards on social media sites and AI chatbots in

*Code and data are available at:

order to protect children from harmful material (Miller 2026). The detection of inappropriate content on websites frequented by minors has thus become a legal necessity, and research in this area is of great interest to social media companies. Seeing as it is virtually impossible for humans to manually check content at this volume, Justice Stewart’s criteria falls short in today’s world.

Instead, companies perform this detection using a mix of algorithms and user reports. In 2021, Meta announced they would begin using artificial intelligence models trained with few-shot prompting to detect harmful content on their platforms (Meta 2021). Few-shot learning is when a model is provided with some small number of completed examples before being given a task (Brown et al. 2020). This is a common method used to help improve the accuracy of LLM text classification. Great strides have been made in recent years to improve the ability of models to classify the semantic meaning of text passages. Text classification is an important area of natural language processing, a subfield of computer science that focuses on getting computers to work with human language. Traditional methods for representing text ignored semantic meaning, and would, for example, treat the word “tie” the same whether it referred to the accessory or an equal score (Altinel and Ganiz 2018). Newer models have overcome this limitation and are able to not only discern between homonyms, but understand tonal differences in text that can be easily picked up on by humans but are very difficult to define. This has vastly broadened the use cases for artificial intelligence.

With the recent advances in semantic text classification technology, much research is being done attempting to use natural language processors as models for human behavior and attitudes. The training corpus for many leading LLMs is essentially the whole internet, containing every human thought ever published online. This creates an interesting opportunity to see how well model responses reflect human populations.

Insert paragraph about methods and results

This rest of this report will be separated into the following sections: Section 2 will provide background information on the key areas most relevant to this paper. Section 3 will showcase the data used in the analysis. Section 4 will describe the methods used [say something else about the methods]. Section 5 will detail the results and findings of the analysis. Section 6 will discuss the big picture relevance of the results.

2 Background

The content of this paper involves three main subjects: (1) childhood sexuality education, (2) online censorship, and (3) semantic text classification. Much research has been done in all three areas that have greatly informed and inspired this report.

2.1 Childhood Sexuality Education

Legislation regarding sex education varies greatly around the world, even within countries such as Canada and the US. An extreme example of this is the United States, which is known for its great political divide. SEICUS, an organization dedicated to promoting healthy sexual education in the US, has created detailed state profiles that go in depth on each region's position on how sex ed should be delivered to students (SEICUS 2025). As of 2025, 14 states provide abstinence-only education, which has been shown to fail in protecting students from sexually transmitted illness and pregnancy (Santelli et al. 2006). 19 states exclude and/or stigmatize LGBTQ+ people in their sex education. Only three states, California; Oregon and Washington, require comprehensive sexuality education in all schools. Different areas have vastly different ideas of what is age-appropriate for those under the age of 18 to learn about. Parents and guardians may have a desire to shelter their children from certain subjects for fear that learning about sexuality and gender will increase risky behaviour, however there is much evidence to the contrary (World Health Organization 2026). A paper on parental attitudes towards the K-12 gender and sexuality curriculum in Australia identified that many believe children are too immature and easily influenced to learn about these topics (Ferfolja, Manlik, and Ullman 2024). Regardless of age-appropriateness, young people may be exposed to sexual concepts through the internet, peers, and real life experiences. Without a reliable foundation of knowledge from which to draw from, they may be unable to tell myth from truth.

2.2 Online Censorship and Age Restriction

As the internet and social media gain increasing popularity, countries have each put in place legislation to regulate what kind of content can be posted online. The Online Streaming Act (Bill C-11) of 2022 gave the Canadian Radio and Television Commission the ability to regulate Canadian digital media platforms the same way they had been regulating radio and TV broadcasters (Woolf 2022). Online censorship is a hotly debated topic, with many feeling it infringes on human rights; some feel the internet should be like the high seas, out of the jurisdiction of any governing force (Howard and Arimateia da Cruz 2022). And in many ways it is, which is part of why it is so difficult to regulate online content. But many social media platforms are getting increasingly strict with their terms of service due to financial and legal pressure (Liao 2018). There is much demand to raise the minimum age requirement for users, however the most effective way to enforce this is by requiring legal identification, which creates privacy issues (Maimann 2025). In 2025, Australia banned social media for those under the age of 16, and several countries, such as Canada, are heading down the same path (2026). The Canadian government has stated that they intend to implement a social media ban for those under the age of 16, with an exemption for platforms that have sufficient safeguards in place (Miller 2026). These safeguards must protect children from viewing “harmful content”, which is defined as media that falls into any of these categories:

- Intimate content communicated without consent

- Content that sexually victimizes a child or revictimizes a survivor
- Content that induces a child to harm themselves
- Content used to bully a child
- Content that foments hatred
- Content that incites violence
- Terrorism or violent extremism content

Unfortunately many of these types of content cannot be fully accurately detected by a simple algorithm. There is no one word that is in every piece of hate speech, no single phrase that consistently promotes violence. Platforms often use user reports to aid algorithms in this detection, however this opens the door to unwarranted mass reporting performed by bots or spiteful users (**citeConteras?**). Content by LGBTQ+ individuals is often disproportionately suppressed, and sexual education is usually detected as explicit (Tremblay 2026) (York 2024) (Zhang and Zou 2025). Additionally, many posts that arguably do fit the criteria go unremoved (Seitz 2022). It would be impossible for a healthy sample of humans to go through every social media post and assess potential harm. It is in the best interest of companies to build more intelligent algorithms that are able to detect nuanced semantics in order to keep online spaces safe and inclusive.

2.3 Text Classification

Text classification is a common task in machine learning in which a model has to assign a category or label to a piece of text. This category could be the topic the text is relating to, the sentiment, the tone, or something else entirely. There are many different techniques used to complete this task.

2.3.1 Supervised Text Classification

The majority of text classification performed today is supervised (Stryker, n.d.). In this method, sets of pre-classified text (which have been labeled by humans) are fed to the model in order to train it. The model uses these examples to get an idea of what goes in each category. It is then given new pieces of text which it must classify according to this same pattern.

2.3.2 Text Clustering

Text clustering is a type of unsupervised text classification, a method where instead of training the model on labeled examples, it finds patterns within the data. This technique essentially puts texts into groups based on semantic similarity. Texts within the same cluster are assumed to relate to the same topic. After this is performed, the clusters can be inspected to determine which topic they each refer to.

2.3.3 Zero-Shot Learning

Zero-shot, one-shot, and few-shot learning are all prompting techniques used to improve the ability of a pre-trained model to perform a specified task. The model is first given some number of completed examples of said task, and then asked to perform it on new data. Few-shot refers to multiple examples, one-shot refers to one, and zero-shot asks the model to execute the task based solely on what it already knows. In the context of text classification, this would look like simply prompting the model to classify some piece of text, usually into a specified set of categories. Zero-shot is often used when there is no labeled data available with which to provide examples. While this method has less accuracy than one- and few-shot learning, it still proves to be quite successful, particularly with larger, more advanced models (Brown et al. 2020).

3 Data

3.1 *It's Perfectly Normal*

The data used in this analysis is a random sample of sentences collected from the most recent edition of the book *It's Perfectly Normal* by Robbie Harris and Michael Emberley, published in 2014 (Harris and Emberley 2014). The first edition of this book was originally published in 1994 and has since been updated three times (Flynn 2019). *It's Perfectly Normal* is a sexual education book intended for children ages 10 and up, and has been praised for its clear language, comprehensive information, and discussion of the LGBTQ+ community (Perring 2005). Despite this, the picture book has also gained widespread controversy for its detailed illustrations and themes which some consider to be inappropriate for the target audience (Hersher 2014). It has frequently topped the American Library Association's list of most challenged books since its original release over two decades ago (American Library Association 2007).

3.2 Cleaning

TBD

4 Methods

5 Results

6 Discussion

References

2026. *Angus Reid Institute*. <https://angusreid.org/social-media-ban-canada-kids/>.
- Altinel, Berna, and Murat Can Ganiz. 2018. "Semantic Text Classification: A Survey of Past and Recent Advances." *Information Processing & Management* 54 (6): 1129–53. <https://doi.org/https://doi.org/10.1016/j.ipm.2018.08.001>.
- American Library Association. 2007. *American Library Association*. <https://www.ala.org/news/8220its-perfectly-normal8221-tops-alas-2005-list-most-challenged-books>.
- Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. "Language Models Are Few-Shot Learners." <https://arxiv.org/abs/2005.14165>.
- Döring, Nicola. 2024. "Social Media and Sex Education." In. https://doi.org/10.1007/978-3-030-95352-2_131-1.
- Ferfolja, Tania, Kate Manlik, and Jacqueline Ullman. 2024. "Parents' Perspectives on Gender and Sexuality Diversity Inclusion in the k-12 Curriculum: Appropriate or Not?" *Sex Education* 24 (5): 632–47. <https://doi.org/10.1080/14681811.2023.2263846>.
- Flynn, Kitty. 2019. "The Horn Book | Happy Anniversary: It's Perfectly Normal." *The Horn Book*. <https://www.hbook.com/story/happy-anniversary-its-perfectly-normal>.
- Harris, Robie, and Michael Emberley. 2014. *It's Perfectly Normal*. 20th anniversary. Candlewick Press.
- Hersher, Rebecca. 2014. "It May Be 'Perfectly Normal,' but It's Also Frequently Banned." *NPR*. <https://www.npr.org/2014/09/21/350366435/it-may-be-perfectly-normal-but-its-also-frequently-banned>.
- Howard, Travis D., and Jose de Arimateia da Cruz. 2022. "Like the Sea, so Cyberspace: A Brief Exploration of Establishing Norms Through a Maritime Lens." *Journal of Advanced Military Studies* 13 (2): 142–53. <https://doi.org/10.21140/mcu.20221302007>.
- Lattman, Peter. 2007. "The Origins of Justice Stewart's 'i Know It When i See It'." *The Wall Street Journal*. <https://www.wsj.com/articles/BL-LB-4558>.
- Liao, Shannon. 2018. "Tumblr Will Ban All Adult Content Starting December 17th." *The Verge*. <https://www.theverge.com/2018/12/3/18123752/tumblr-adult-content-porn-ban-date-explicit-changes-why-safe-mode>.
- Maimann, Kevin. 2025. "'Can i See Some Id?' As Online Age Verification Spreads, so Do Privacy Concerns." *CBC*. <https://www.cbc.ca/news/online-safety-act-privacy-1.7598113>.

- Meta. 2021. “Harmful Content Can Evolve Quickly. Our New AI System Adapts to Tackle It.” *Meta AI*. <https://ai.meta.com/blog/harmful-content-can-evolve-quickly-our-new-ai-system-adapts-to-tackle-it/>.
- Miller, Marc. 2026. *Safe Social Media Act*. <https://www.canada.ca/en/canadian-heritage/services/safe-social-media-act.html>.
- Perring, Christian. 2005. “Review of ”It’s Perfectly Normal”” *Northern Wyoming Mental Health Center*. https://wyomentalhealth.org/poc/view_doc.php?type=book&id=2483&cn=28.
- Santelli, John, Mary A. Ott, Maureen Lyon, Jennifer Rogers, Daniel Summers, and Rebecca Schleifer. 2006. “Abstinence and Abstinence-Only Education: A Review of u.s. Policies and Programs.” *Journal of Adolescent Health* 38 (1): 72–81. <https://doi.org/https://doi.org/10.1016/j.jadohealth.2005.10.006>.
- SEICUS. 2025. “Siecus State Profiles.” *SEICUS*. <https://siecus.org/siecus-state-profiles/>.
- Seitz, Amanda. 2022. “Social Media Platforms Fail to Stop LGBTQ Hate Speech, According to GLAAD Report.” *PBS*. <https://www.pbs.org/newshour/nation/social-media-platforms-fail-to-stop-lgbtq-hate-speech-according-to-glaad-report>.
- Stryker, Cole. n.d. “What Is Text Classification?” *IBM*. <https://www.ibm.com/think/topics/text-classification>.
- Tremblay, Chris. 2026. “Transgender Creators Confront ‘Pink Shadow Ban’ as TikTok and Meta Face Scrutiny over Content Suppression.” *Edge*. <https://www.edgemedianetwork.com/story/168090/transgender-creators-confront-pink-shadow-ban-as-tiktok-and-meta-face-scrutiny>.
- Woolf, Marie. 2022. “MPs Rush over 150 Amendments to Streaming Bill, Senate Says It Won’t Be Pressured | CBC News.” *CBC*. <https://www.cbc.ca/news/politics/streaming-bill-committee-amendments-1.6490285>.
- World Health Organization. 2026. “Comprehensive Sexuality Education.” *World Health Organization*. <https://www.who.int/news-room/fact-sheets/detail/comprehensive-sexuality-education>.
- York, Jillian. 2024. “The Global Suppression of Online LGBTQ+ Speech Continues.” *EFF*. <https://www.eff.org/deeplinks/2024/06/global-suppression-online-lgbtq-speech-continues>.
- Zhang, Xinyu, and Wenxue Zou. 2025. “Shattering the Silence: Exploring the Role of Chinese Online Sexual Health Influencers in Promoting Sex Education.” *Sex Education* 25 (1): 95–110. <https://doi.org/10.1080/14681811.2023.2272734>.