

Classifying Sexual Education Material with Generative AI*

Understanding how generative models classify the age appropriateness of sentence excerpts from sexual education material

Rosemarie Topp

August 25, 2026

While sexual education is a critical part of child development, there is debate around how and when it should be taught. Social media companies use large language models tools to flag content that is inappropriate for younger audiences. Using zero-shot learning, this report examines how generative AI classifies the age appropriateness of sexual education material in order to understand what language factors into this decision. Model performance is then evaluated based on how it compares to the general model consensus.

1 Introduction

Due to the subject matter, this paper contains sensitive language that may be considered offensive.

The state of sexuality education in schools varies around the world. The topic of if, when, and how to introduce children to certain subjects is divisive, and often raises religious, political, and moral issues. Comprehensive sexuality education (CSE) is “a curriculum-based, incremental and scientifically-accurate approach that teaches children and young people about the cognitive, emotional, physical and social aspects of sexuality in a culturally-relevant and age-appropriate way.” (World Health Organization 2026). That being said, what is considered age-appropriate for a given demographic is hard to define. Even if it could be universally agreed upon that a piece of content was inappropriate for children to consume, the exact reasons why may be hard to put into words. The exact criteria often comes down to what United States Supreme Court Justice Potter Stewart once said when asked to define obscene pornography: “I know it when I see it” (Lattman 2007).

*Code and data are available at: https://github.com/rosiet1039/sex_ed_classification_paper

In 2024, 66,000 Instagram posts, 347,200 Tweets, and 1.7 million Facebook posts were shared every minute (Domo 2024). Some pieces of content may have contained sensitive themes. Websites such as Youtube or TikTok are frequented by children seeking educational information (Döring 2024). In the early 2000s, before censorship laws were put in place, the internet was essentially unregulated, allowing anyone with internet access to consume and post offensive, violent, and pornographic material. Over the years, various countries have put legislation in place to help hide sensitive content from younger users, such as Canada’s recent Bill C-34, which mandates the implementation of safeguards on social media sites and AI chatbots in order to protect children from harmful material (Miller 2026). The detection of inappropriate content on websites frequented by minors has thus become a legal necessity, making research in the area of data classification of great interest to social media companies. Seeing as it is virtually impossible for humans to manually check content at this volume, Justice Stewart’s criteria falls short in today’s world.

Instead, companies perform this detection using a mix of algorithms and user reports. In 2021, Meta announced they would begin using artificial intelligence models trained with few-shot learning to detect harmful content on their platforms (Meta 2021). Few-shot learning is a technique in which a model is provided with some small number of demonstrative examples before a task is given (Brown et al. 2020). This is a common method used to help improve the accuracy of LLM text classification. Text classification is an important area of natural language processing, a subfield of computer science that focuses on getting computers to work with human language. Traditional methods for representing text ignored semantic meaning, and would, for example, treat the word “tie” the same whether it referred to the accessory or an equal score (Altinel and Ganiz 2018). Newer models have overcome this limitation and are able to not only discern between homonyms, but understand tonal differences in text that can be easily picked up on by humans but are very difficult to define. This has vastly broadened the use cases for artificial intelligence.

With the recent advances in semantic text classification technology, much research is being done attempting to use natural language processors as models for human behavior and attitudes. The training corpus for many leading LLMs is essentially the whole internet, containing every human thought ever published online. This creates an interesting opportunity to see how well model responses reflect human populations.

It’s Perfectly Normal is a ten-and-up sexual education book by Robie Harris and Michael Emberley (Harris and Emberley 2014). Though it has been praised for its detail and comprehensiveness, the book has gained criticism from those who argue that its contents are not appropriate for the target demographic (Perring 2005) (Hersher 2014). A random sample of 100 sentences was taken from the book, and using zero-shot learning, five different models were prompted to classify whether or not each text excerpt is appropriate for ages ten and up.

Model responses were then examined using scatterplots of two dimensional embeddings of each sentence. Sentences that discussed topics such as arousal, puberty, external genitalia, and sexual health tended to be flagged as inappropriate most often by all models. Performances were evaluated using alternative recall and precision metrics which compare individual classifications

to overall consensus. It was found that GPT 5.4 Mini had the highest recall (79.0%) and the lowest precision (35.4%), while Claude Haiku had the inverse, (with scores 21.3% and 51.3% and respectively). GPT 4o had the best F score (49.3).

This rest of this report will be separated into the following sections: Section 2 will provide background information on the key areas most relevant to this paper. Section 3 will go through relevant literature and related analyses. Section 4 will showcase the data used in the analysis. Section 5 will describe the models and prompting technique. Section 6 will detail the results and findings of the analysis. Section 7 will discuss the big picture relevance of the results.

2 Background

The content of this paper involves three main subjects: (1) childhood sexuality education, (2) online censorship, and (3) semantic text classification. Much research has been done in all three areas that have greatly informed and inspired this report.

2.1 Childhood Sexuality Education

Legislation regarding sex education varies greatly around the world, even within countries such as Canada and the United States. An extreme example of this is the United States, which is known for its great political divide. SEICUS, an organization dedicated to promoting healthy sexual education in the US, has created detailed state profiles that go in depth on each region's position on how sex ed should be delivered to students (SEICUS 2025). As of 2025, 14 states provide abstinence-only education, which has been shown to fail in protecting students from sexually transmitted illness and pregnancy (Santelli et al. 2006). 19 states exclude and/or stigmatize LGBTQ+ people in their sex education. Only three states, California; Oregon and Washington, require CSE in all schools.

Different areas have vastly different ideas of what is age-appropriate for those under the age of 18 to learn about. Parents and guardians may have a desire to protect their children from certain subjects for fear that learning about sexuality and gender will increase risky behaviour, however there is much evidence to the contrary (World Health Organization 2026). A paper on parental attitudes towards the K-12 gender and sexuality curriculum in Australia identified that many believe children are too immature and easily influenced to learn about these topics (Ferfolja, Manlik, and Ullman 2024). Regardless of age-appropriateness, young people may be exposed to sexual concepts through the internet, peers, and real life experiences. Without a reliable foundation of knowledge from which to draw from, they may be unable to tell myth from truth.

2.2 Online Censorship and Age Restriction

As the internet and social media gain increasing popularity, countries have each put in place legislation to regulate what kind of content can be posted online. The Online Streaming Act (Bill C-11) of 2022 gave the Canadian Radio and Television Commission the ability to regulate Canadian digital media platforms the same way they had been regulating radio and TV broadcasters (Woolf 2022). This bill received pushback from those who hold the opinion that online censorship infringes on human rights (Lewis 2023). Regardless, many social media platforms are getting increasingly strict with their terms of service due to financial and legal pressure, such as content platform OnlyFans banning sexually explicit content at the request of credit card companies (CNN 2021) (Liao 2018). The Canadian government has stated that they intend to implement a social media ban for those under the age of 16, with an exemption for platforms that have sufficient safeguards in place (Miller 2026). These safeguards must protect children from viewing “harmful content”, which is defined as media that falls into any of these categories:

- Intimate content communicated without consent
- Content that sexually victimizes a child or revictimizes a survivor
- Content that induces a child to harm themselves
- Content used to bully a child
- Content that foments hatred
- Content that incites violence
- Terrorism or violent extremism content

Unfortunately many of these types of content cannot be fully accurately detected by a simple algorithm. There is no one word that is in every piece of hate speech, no single phrase that consistently promotes violence. Platforms often use user reports to aid algorithms in this detection, however this opens the door to unwarranted mass reporting performed by bots or spiteful users (Contreras 2021). Content by LGBTQ+ individuals is often disproportionately suppressed, and sexual education is usually detected as explicit (Tremblay 2026) (York 2024) (Zhang and Zou 2025). Additionally, many posts that arguably do fit the criteria go unremoved (Seitz 2022). It would be impossible for a healthy sample of humans to go through every social media post and assess potential harm. It is in the best interest of companies to build more intelligent algorithms that are able to detect nuanced semantics in order to keep online spaces safe and inclusive.

2.3 Text Classification

Text classification is a common task in machine learning in which a model has to assign a category or label to a piece of text. This category could be the topic the text is relating to, the sentiment, the tone, or something else entirely. There are many different techniques used to complete this task.

The majority of text classification performed today is supervised (Stryker, n.d.). In this method, sets of pre-classified text (which have been labeled by humans) are fed to the model in order to train it. The model uses these examples to get an idea of what goes in each category. It is then given new pieces of text which it must classify according to this same pattern.

Text clustering is a type of unsupervised text classification, a method where instead of training the model on labeled examples, it finds patterns within the data. This technique essentially puts texts into groups based on semantic similarity. Texts within the same cluster are assumed to relate to the same topic. After this is performed, the clusters can be inspected to determine which topic they each refer to.

Zero-shot, one-shot, and few-shot learning are all prompting techniques used to improve the ability of a pre-trained model to perform a specified task. The model is first given some number of completed examples of said task, and then asked to perform it on new data. Few-shot refers to multiple examples, one-shot refers to one, and zero-shot asks the model to execute the task based solely on what it already knows. In the context of text classification, this would look like simply prompting the model to classify some piece of text, usually into a specified set of categories. Zero-shot is often used when there is no labeled data available with which to provide examples. While this method has less accuracy than one- and few-shot learning, it still proves to be quite successful, with GPT-3 achieving 64.3% accuracy on TriviaQA, a dataset of question and answer problems, and 76% on LAMBADA, a dataset used to test word prediction capabilities (Brown et al. 2020).

3 Related Works

The ability of LLMs to reflect human attitudes has been explored across several fields. Argyle et al. (2023) conditions GPT-3 on thousands of socio-demographic backstories from United States surveys in order to create “silicon samples” of human beings, and analyzes the model’s ability to mirror the attitudes, behaviours, and opinions of said demographics. The results suggest that with sufficient conditioning, GPT-3 can be a novel tool for studying human populations in the absence of real subjects. Horton, Filippas, and Manning (2026) similarly endows several models with various personas in order to analyze how LLMs behave in simulated economic experiments. They also find that LLMs can be used as a fast and inexpensive way to perform preliminary economic research.

A useful application of this simulation of human attitudes is to give LLMs the task of interpreting the semantic meaning of text and other media. Galke et al. (2025) shows that pre-trained models like GPT, while outperformed by encoder-only pre-trained models, perform text classification more accurately than many newer methods. Chiu, Collins, and Alexander (2022) uses zero-, one-, and few-shot learning to analyze GPT-3’s ability to classify text as racist, sexist, both, or neither. It is found to have an average accuracy of 55% to 67%, and accuracy up to 85% in a few-shot setting. This paper provides convincing evidence that LLMs such as GPT-3, with the help of few-shot learning, can be useful in detecting harmful content online.

George and Surdeanu (2023) introduces a dataset of content that has been labeled sexually suggestive, sexually educative, or neither, from the social media platform TikTok. This labeled data is useful for training LLMs to improve detection of explicit content while avoiding the censorship of educational material. They validate the usefulness of this dataset by training two LLMs to classify its contents, and find that while difficult, the task is learnable.

4 Data

4.1 Overview

The data analyzed in this paper comes from the 2014 edition of the book *It's Perfectly Normal* by Robie Harris and Michael Emberley. All handling of data was done using the statistical programming language R (R Core Team 2025) as well as Python version 3.13 (Python Software Foundation 2026). R packages `readr` (Wickham, Hester, and Bryan 2024), `text` (Kjell, Giorgi, and Schwartz 2023), `tidyverse` (Wickham et al. 2019), and `uwot` (Melville 2025), and Python packages `pymupdf` (McKie and Giese 2026), `nltk` (Bird and Loper 2004), `re` (Foundation 2026), and `pandas` (Team 2020) were used to clean and analyze data. R package `ellmer` (Wickham et al. 2026) was used to access the models used in this analysis. All plots were created using `ggplot2` (Wickham 2016), `ggrepel` (Slowikowski 2024), and `ggpubr` (Kassambara 2026). `knitr` (Xie 2014) and `kableExtra` (Zhu 2024) were used to help format this document.

4.2 *It's Perfectly Normal*

The data used in this analysis is a random sample of sentences collected from the most recent edition of the book *It's Perfectly Normal* by Robie Harris and Michael Emberley, published in 2014 (Harris and Emberley 2014). The first edition of this book was originally published in 1994 and has since been updated three times (Flynn 2019). *It's Perfectly Normal* is a sexual education book intended for children ages 10 and up, and has been praised for its clear language, comprehensive information, and discussion of the LGBTQ+ community (Perring 2005). Despite this, the picture book has also gained widespread controversy for its detailed illustrations and themes which some consider to be inappropriate for the target audience (Hersher 2014). It has frequently been in the top 10 of the American Library Association's list of most challenged books since its original release over two decades ago (American Library Association 2007).

4.3 Cleaning

The 20th anniversary edition of *It's Perfectly Normal* is 116 pages front to back including the cover, and contains 30 chapters, including a one page introduction. In order to analyze the book, all text was extracted from a pdf of the book and parsed into sentences, which were then written into a csv file. Only text with font size 10.5 was included, which corresponded

to the main body text. This resulted in a dataset of 1,214 sentences, spanning a total of 85 pages and 29 chapters, including the introduction (note that chapter 6 was not included as it contained no main body text). Summary statistics of the included number of sentences per chapter can be seen in **?@fig-perchap**. While effort was taken to preserve the original format, the text parsing software only included the first segment of sentences that spanned multiple pages. These fragments were left in the dataset in order to avoid loss of data. Three of these sentences ended up in the random sample referenced in Section 5, and these were manually completed by referencing the original text. An example of a manually completed sentence is below:

Incomplete sentence from page 62: Sometimes, if a female becomes very sick, she may have

Manually completed sentence from pages 62 and 63: Sometimes, if a female becomes very sick, she may have her eggs removed by a doctor before she takes medication.

Table 1: Distribution of Sentences*

	Mean	S.D.
Per Chapter	41.86207	25.42598
Per Page	14.28235	6.99083

*Over included chapters and pages

5 Methods

Several models were used to examine how generative LLMs classify text-based sexual education material. Default parameters were maintained on all models except for temperature. The temperature of a model indicates the variability of its responses, with 0 being the least variable and 1 the most. This parameter was set to 0.4 in order to examine how model response varies. The models used were Open AI’s GPT 4o mini and GPT 5.4 mini, Deepseek’s V4 Flash and V4 Pro, and Claude Haiku 4.5 (OpenAI et al. 2024) (Singh et al. 2026) (DeepSeek-AI 2026) (Anthropic 2025).

In order to prompt the models, zero-shot learning was employed. Zero-shot refers to a prompting style where no demonstrative examples are provided, and the model must base its answer on previous knowledge. A random sample of 100 unique sentence excerpts from *It’s Perfectly Normal* were given to each model 30 times. The prompt template is as follows: ‘Is the following text in quotes appropriate for children ages 10 and up? Answer yes or no [sentence in quotations]’. The only deviation to this format was in the case of Claude Haiku 4.5, which disregarded the instruction to respond with yes or no. The model did, however, comply when ‘Answer yes or no’ was changed to ‘Answer with only yes or no’, and thus this slight prompt variation was used.

6 Results

Models all gave yes or no answers to each prompt. All response variations can be found in **?@fig-responses**.

Note that here on out, a “positive” classification refers to a negative model response, or in other words a classification of inappropriate.

Table 2: Model Responses

model	responses
GPT 4o Mini	Yes., No., Yes
GPT 5.4 Mini	Yes., No., Yes, No
Deepseek V4 Flash	Yes., Yes, No, No., yes, no
Deepseek V4 Pro	Yes, yes, no, No
Claude Haiku 4.5	Yes, No, No., Yes

6.1 Semantic Embedding and Clustering

In order to understand what kind of language tended to be classified as inappropriate, semantic text embeddings were created for each sentence. This refers to the process of translating text into a long list of numbers called a vector, which points to an area in multi-dimensional space. This translation takes into account semantic meaning, rather than just the words themselves, so sentences that mean similar things will be located closer together. This was done using function ‘textEmbed’ from the ‘text’ package in R, with model bert-case-uncased, which can be found for free on HuggingFace (Kjell, Giorgi, and Schwartz 2023) (Devlin et al. 2018).

The sentences were then divided into clusters using the Hartigan and Wong (1979) K-means algorithm. This algorithm attempts to find the optimal way to divide data into some pre specified k groups called clusters. This grouping is done based on the Euclidean distance (straight length) between points. In this case, clustering helps to identify the more specific themes at play within the sample of sentences. K=3 was chosen for the number of clusters, and the following themes were identified:

- 1: Myths, Life Advice, and Internet Safety
- 2: Healthcare, STIs, Arousal, Puberty, and External Anatomy
- 3: Internal Biology and Anatomy

The percentage of sentences positively classified in each grouping can be found in **?@fig-bycluster**. These are some examples of sentences in each cluster:

Cluster 1:

Even though that person is not actually touching your body, talking about sex or your body in this way can be a kind of sexual abuse.

If you say online that someone is fat or skinny or sexy or ugly or beautiful or handsome, what you have said is really never private once those words are on the Internet.

Some websites are not always what they tell you they are.

Cluster 2:

The IUD, the implant, and Depo-Provera are the most effective kinds of birth control.

Others believe that using the rhythm method and withdrawal method is fine.

And yet people, both kids and adults, who are HIV positive or who have developed AIDS have been discriminated against.

Cluster 3:

If you are male, your penis and testicles are two of your sex organs.

Identical twins begin if a single egg splits into two after it has been fertilized.

In most births, the baby’s head pushes out of the vagina first.

It was decided that three clusters formed the most clear thematic division, however further insight could potentially be gained by grouping the data into some other number of clusters. Additionally, the above themes were identified by manually inspecting the sentences in each cluster, but other terms could be utilized. An alternative method for labeling clusters is to feed a sample of sentences from each group to a generative model and prompting it to identify themes.

The embeddings were then compressed into two dimensions using Uniform Manifold Approximation and Projection (UMAP), a dimension reduction technique, with cosine similarity (angular distance between vectors) as the chosen distance metric. This was done in order to visualize the embeddings in the form of scatterplots in Figure 1.

Table 3: Percentage Positively Classified by Cluster

Cluster	Model	Mean (%)	Standard Deviation (%)
	GPT 4o Mini	0.0	0.0
	GPT 5.4 Mini	9.3	1.3

1	Deepseek V4 Flash	0.6	1.2
	Deepseek V4 Pro	6.8	2.5
	Claude Haiku 4.5	0.0	0.0
2	GPT 4o Mini	17.4	1.8
	GPT 5.4 Mini	24.5	2.8
	Deepseek V4 Flash	10.9	2.9
	Deepseek V4 Pro	10.7	4.0
	Claude Haiku 4.5	7.7	0.0
	GPT 4o Mini	3.7	3.3
3	GPT 5.4 Mini	19.9	4.2
	Deepseek V4 Flash	0.9	1.7
	Deepseek V4 Pro	1.4	1.9
	Claude Haiku 4.5	0.0	0.0

6.2 Performance Metrics

Due to a lack of reliable labelled data, model responses could not be compared to a ground truth. Model performances were thus measured using alternative precision and recall metrics found in Fedorchuk and Lamiroy (2017), which compare individual classifications to overall model consensus. Consider the probability $P(\delta_i)$ of sentence δ_i to be positively classified. Then:

$$P(\delta_i) = \sum_{k=1}^s \frac{S_k(\delta_i)}{s}$$

Where

$$S_k(\delta_i) = \begin{cases} 0 & \text{if } Negative \\ 1 & \text{if } Positive \end{cases}$$

is the classification result from model S_k of sentence δ_i , and s is the total number of models

Assuming sentences are evenly distributed, valid probabilistic measures of precision $Pr(S_k)$ and recall $Rc(S_k)$ are as follows:

$$Pr(S_k) = \sum_{i=1}^d \frac{S_k(\delta_i)P(\delta_i)}{S_k(\delta_i)} Rc(S_k) = \sum_{i=1}^d \frac{S_k(\delta_i)P(\delta_i)}{P(\delta_i)}$$




Figure 1: Clustered Sentences Embeddings

Then the F score is the harmonic mean of these two metrics:

$$F(S_k) = \frac{2Pr(S_k)Rc(S_k)}{Pr(S_k) + Rc(S_k)}$$

These statistics were calculated for each model 30 times across the 30 trials. The (arithmetic) mean and standard deviation of each metric are shown in [?@fig-metrics](#).

Table 4: Performance Evaluation Metrics by Model

Model	Metric	Mean (%)	Standard Deviation (%)
GPT 4o Mini	Recall	48.4	5.5
	Precision	50.6	4.5
	F Score	49.3	4.0
GPT 5.4 Mini	Recall	79.0	6.1
	Precision	35.4	2.3
	F Score	48.9	2.8
Deepseek V4 Flash	Recall	30.9	8.5
	Precision	53.9	6.4
	F Score	38.6	8.3
Deepseek V4 Pro	Recall	38.3	10.9
	Precision	45.2	6.8
	F Score	40.8	8.2
Claude Haiku 4.5	Recall	21.3	2.6
	Precision	57.3	7.9
	F Score	31.1	3.7

7 Discussion

As the sentence embedding scatterplots demonstrate, models all tended to classify topics as inappropriate in the second cluster (highlighted in green) more than the other two. Manual inspection of the texts in this cluster showed they pertain to sexual health, arousal, puberty, and external anatomy. In order to understand how the models were classifying certain specific topics, sentences relating to avoiding pregnancy, sexually transmitted infections, arousal, puberty, females, and males were manually picked. Aggregated statistics on how these sentences were classified can be found in Figure 2. It should be noted that alternative sentence groupings could

be chosen, for example the sentence “Sex hormones can also affect girls’ and boys’ feelings and moods.” could have been included under the topic of puberty, but was not in this analysis.

GPT 4o Mini and GPT 5.4 Mini tend to exhibit similar trends in terms of which topics are most and least appropriate. The two Deepseek models also have similarities.

All models agree that puberty is an appropriate subject for children ten and up to learn about, which aligns with the biological reality that puberty begins around 8 for girls and 9 for boys (Clinic 2024).

The models classified 10%-40% of sentences about avoiding pregnancy positively. In Ontario, it is suggested to teach students about “how to (...) avoid becoming a parent before they are ready” in the seventh grade, which is when most students are aged 12 to 13 (Ontario 2026). The *International Technical Guidance on Sexuality Education* (ITGoSE), however, recommends children start learning about contraceptives as young as 9, but suggests teaching students about abortion around the age of 15 to 18 or older (UNESCO 2018). Sentences about STIs were classified positively 5%-25% of the time, which was lower than the previous topic by around 10 to 20 percentage points for all models except GPT 4o Mini. This is surprising as the two subjects tend to be taught together; both the province of Ontario and ITGoSE recommend introducing the subject at the same time as conversations about avoiding pregnancy (Ontario 2026) (UNESCO 2018). This could be related to online anti-choice sentiments in training datasets that may discourage discussions pertaining to pregnancy avoidance and termination (Tavrow P 2025).

Arousal was considered the most inappropriate topic by both GPT models, with GPT 4o Mini classifying 55.0% and GPT 5.4 Mini classifying 50.83% positively. The Deepseek models only considered 10%-13% of these sentences to be inappropriate, and Claude considered this topic entirely appropriate, a noticeable difference between models. The province of Ontario mentions teaching 6th graders (around 11 to 12 years of age) about the adverse effects of consuming sexually explicit media, but otherwise does not mention the subject. ITGoSE explicitly names arousal, suggesting it be brought up around the ages of 9 to 12. Reproduction, however, is one of the first sexual education lessons taught to children; it is suggested in ITGoSE to teach students as young as 5 about the process of reproduction. Each model classified reproduction as more appropriate than arousal, yet arousal is an essential part of the reproductive process. Many cultures cultivate stigma around pleasure, especially when it is not required for insemination, while sanctifying reproduction (Verena Klein 2024) (Bryan Kutner 2025). A clear example of this is in the case of the Virgin Mary, whose immaculate conception required no sexual contact whatsoever (Tinsley 2005). It has also been found in several studies that a significant portion of parents (43% in Hurst et al. (2024), 45% in Venka Simovska et al. (2015)) do not want their children to learn about the pleasurable side of sex, potentially due to fear that it would incentivize the act (Pariera 2016). The discrepancy in classification of sentences related to reproduction versus arousal indicates that models are absorbing anti-pleasure attitudes within their training corpora.

Another notable difference between models is the classification of male and female topics. Both GPT 4o Mini and GPT 5.4 Mini positively classified sentences about men with around 20 to 30 percentage points more than that relating to women, while the remaining three models all displayed the inverse behaviour, though there was less of a difference in percentages. This could be due to the overlap between sentences about women and sentences about avoiding pregnancy. The topic of pregnancy avoidance placed near the middle in terms of appropriateness for both GPT models, while it was considered the least appropriate topic for the other three, and thus this could explain the discrepancy. Further research examining gendered classification differences on more controlled data would provide deeper insight.

Table 5: Percentage of Sentences Classified as Inappropriate by Topic

Topics	GPT 4o (%)	GPT 5.4 (%)	Deepseek Flash (%)	Deepseek Pro (%)	Claude Haiku (%)
Avoiding Pregnancy	11.67	38.00	27.67	28.00	20.00
STIs	14.85	24.85	6.36	9.09	9.09
Arousal	55.00	50.83	12.92	10.00	0.00
Puberty	0.00	0.37	0.00	0.00	0.00
Reproduction	5.67	22.67	0.00	0.00	0.00
Female	8.33	16.67	12.22	10.42	8.33
Male	28.33	46.39	0.83	3.89	0.00

Figure 2

Due to a lack of ground truth data, model performances were measured using recall and precision metrics that compare responses to the overall model consensus. According to these statistics, GPT 5.4 had the highest mean recall (79.0%) and lowest mean precision (35.4%), meaning it correctly classified the most inappropriate sentences, but incorrectly classified the most appropriate ones. Claude Haiku demonstrated the opposite behaviour, with the lowest mean recall (21.3%) and highest mean precision (57.3%). This model was much more conservative when classifying text as inappropriate, only putting three sentences in this category out of 100 prompts and 30 trials for each prompt. It should be noted that Claude Haiku had no variance in its responses, and always responded with Yes or No for each unique prompt. GPT 4o Mini had the highest mean F score (49.3%), demonstrating the most balanced behaviour overall, while Claude Haiku had the lowest (31.1%).

All of these scores fell within the range of 30% to 60%, with the exception of GPT 5.4 Mini’s recall score of 79.0%. Rather than indicating poor performance across all models, as they would if these metrics were based on ground truth, they demonstrate a lack of model consensus on a threshold for content age appropriateness. This aligns with research on parental attitudes around sexual education, which found variability even among those who believe in CSE should be taught in schools.

In conclusion, models have varying thresholds for what content can be considered age appropriate for children ten and older, both in terms of particular topics, as well as the overall proportion of positive classifications. Similarities in topic classification were found between different versions

of the same model, and the two Deepseek models somewhat aligned with Claude Haiku 4.5. All models tended to positively classify sentences regarding arousal, stis, and avoiding pregnancy more than other topics, with arousal being the most positively classified for the two GPT models, and avoiding pregnancy for the other three. Discrepancies were found between models regarding classification of sentences pertaining to men versus women, which may be due to topic overlap. Recall, precision, and F measure metrics found that GPT 5.4 Mini had the highest recall but lowest precision, Claude Haiku 4.5 had the highest precision and lowest recall, and GPT 4o Mini had the overall highest F measure, indicating it may be the best model for this task.

A major limitation to this study is the lack of ground truth data with which to compare model responses. If reliable human-labeled data regarding attitudes around sexual education material were to become available, this would be an interesting area of investigation, and would allow for much better model evaluation. It would also be very valuable to replicate this research on other sexual education material, such as that found in another book or on the internet. Additionally, investigation into model classification of non-educational sexual material would be valuable to gain broader insight into model attitudes on sexuality as a whole. Research might also be done on a wider set of models, as this would also allow comparison between model size and classification. Some other variables that could be changed are model parameters such as temperature, how sentences were grouped into topics, and prompting style. The age threshold within the prompt was chosen based on the age recommendation in *It's Perfectly Normal*, however insight could be gained by looking into how classification changes based on the age group at hand, and these results could be compared to sexual education guidelines such as ITGoSE.

References

- Altinel, Berna, and Murat Can Ganiz. 2018. "Semantic Text Classification: A Survey of Past and Recent Advances." *Information Processing & Management* 54 (6): 1129–53. <https://doi.org/https://doi.org/10.1016/j.ipm.2018.08.001>.
- American Library Association. 2007. *American Library Association*. <https://www.ala.org/news/8220its-perfectly-normal8221-tops-alas-2005-list-most-challenged-books>.
- Anthropic. 2025. "System Card: Claude Haiku 4.5." <https://www-cdn.anthropic.com/963373e433e489a87a10c823c52a0a013e9172dd.pdf>.
- Argyle, Lisa P., Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. 2023. "Out of One, Many: Using Language Models to Simulate Human Samples." *Political Analysis* 31 (3): 337–51. <https://doi.org/10.1017/pan.2023.2>.
- Bird, Steven, and Edward Loper. 2004. "NLTK: The Natural Language Toolkit." In *Proceedings of the ACL Interactive Poster and Demonstration Sessions*, 214–17. Barcelona, Spain: Association for Computational Linguistics. <https://aclanthology.org/P04-3031/>.

- Brown, Tom B., Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. “Language Models Are Few-Shot Learners.” <https://arxiv.org/abs/2005.14165>.
- Bryan Kutner, Yumeng Wu et al., Kathrine Meyers. 2025. “Mechanisms of Action and Preliminary Impact of a Stigma-Mitigation Training to Enable Healthcare Discussions about Anal Pleasure and Health in China: A Sequential Explanatory Mixed Methods Pilot Study.” <https://rdcu.be/fBEay>.
- Chiu, Ke-Li, Annie Collins, and Rohan Alexander. 2022. “Detecting Hate Speech with GPT-3.” <https://arxiv.org/abs/2103.12407>.
- Clinic, Cleveland. 2024. “Puberty.” <https://my.clevelandclinic.org/health/body/puberty>.
- CNN. 2021. “Why Did OnlyFans Ban Sexually Explicit Content? It Says It’s the Credit Card Companies.” <https://www.ctvnews.ca/sci-tech/article/why-did-onlyfans-ban-sexually-explicit-content-it-says-its-the-credit-card-companies/>.
- Contreras, Brian. 2021. “‘I Need My Girlfriend Off TikTok’: How Hackers Game Abuse-Reporting Systems.” *LA Times*. <https://www.latimes.com/business/technology/story/2021-12-03/inside-tiktoks-mass-reporting-problem>.
- DeepSeek-AI. 2026. “DeepSeek-V4: Towards Highly Efficient Million-Token Context Intelligence.” <https://arxiv.org/abs/2606.19348>.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. “BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding.” *CoRR* abs/1810.04805. <http://arxiv.org/abs/1810.04805>.
- Domo. 2024. “Data Never Sleeps 10.0.” <https://www.domo.com/data-never-sleeps>.
- Döring, Nicola. 2024. “Social Media and Sex Education.” https://doi.org/10.1007/978-3-030-95352-2_131-1.
- Fedorchuk, Maksym, and Bart Lamiroy. 2017. “Binary Classifier Evaluation Without Ground Truth.” In *2017 Ninth International Conference on Advances in Pattern Recognition (ICAPR)*, 1–6. <https://doi.org/10.1109/ICAPR.2017.8593175>.
- Ferfolja, Tania, Kate Manlik, and Jacqueline Ullman. 2024. “Parents’ Perspectives on Gender and Sexuality Diversity Inclusion in the k-12 Curriculum: Appropriate or Not?” *Sex Education* 24 (5): 632–47. <https://doi.org/10.1080/14681811.2023.2263846>.
- Flynn, Kitty. 2019. “The Horn Book | Happy Anniversary: It’s Perfectly Normal.” *The Horn Book*. <https://www.hbook.com/story/happy-anniversary-its-perfectly-normal>.
- Foundation, Python Software. 2026. *Python Language Reference, Version 3.14*. <https://docs.python.org/3/library/re.html>.
- Galke, Lukas, Ansgar Scherp, Andor Diera, Fabian Karl, Bao Xin Lin, Bhakti Khera, Tim Meuser, and Tushar Singhal. 2025. “Are We Really Making Much Progress in Text Classification? A Comparative Review.” <https://arxiv.org/abs/2204.03954>.
- George, Enfa, and Mihai Surdeanu. 2023. “It Is Not Sexually Suggestive, It Is Educational. Separating Sex Education from Suggestive Content on TikTok Videos.” <https://arxiv.org/abs/2307.03274>.
- Harris, Robie, and Michael Emberley. 2014. *It’s Perfectly Normal*. 20th Anniversary. Candlewick Press.
- Hartigan, J. A., and M. A. Wong. 1979. “Algorithm AS 136: A k-Means Clustering Algorithm.”

- Journal of the Royal Statistical Society. Series C (Applied Statistics)* 28 (1): 100–108. <http://www.jstor.org/stable/2346830>.
- Hersher, Rebecca. 2014. “It May Be ‘Perfectly Normal,’ but It’s Also Frequently Banned.” *NPR*. <https://www.npr.org/2014/09/21/350366435/it-may-be-perfectly-normal-but-its-also-frequently-banned>.
- Horton, John J., Apostolos Filippas, and Benjamin S. Manning. 2026. “Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus?” <https://arxiv.org/abs/2301.07543>.
- Hurst, Jeffrey L., Laura Widman, Julia Brasileiro, Anne J. Maheux, Reina Evans-Paulson, and Sophia Choukas-Bradley. 2024. “Parents’ Attitudes Towards the Content of Sex Education in the USA: Associations with Religiosity and Political Orientation.” *Sex Education* 24 (1): 108–24. <https://doi.org/10.1080/14681811.2022.2162871>.
- Kassambara, Alboukadel. 2026. *Ggpubr: ‘Ggplot2’ Based Publication Ready Plots*. <https://CRAN.R-project.org/package=ggpubr>.
- Kjell, Oscar, Salvatore Giorgi, and H. Andrew Schwartz. 2023. “The Text-Package: An r-Package for Analyzing and Visualizing Human Language Using Natural Language Processing and Deep Learning.” *Psychological Methods*. <https://doi.org/10.1037/met0000542>.
- Lattman, Peter. 2007. “The Origins of Justice Stewart’s ‘i Know It When i See It.’” *The Wall Street Journal*. <https://www.wsj.com/articles/BL-LB-4558>.
- Lewis, Leslyn. 2023. “Debating Bill c-11—Internet Censorship for Canadians.” YouTube. March 31, 2023. <https://www.youtube.com/watch?v=bmd4sASKFRw>.
- Liao, Shannon. 2018. “Tumblr Will Ban All Adult Content Starting December 17th.” *The Verge*. <https://www.theverge.com/2018/12/3/18123752/tumblr-adult-content-porn-ban-date-explicit-changes-why-safe-mode>.
- McKie, Jorj X., and Christian Giese. 2026. “PyMuPDF: High-Performance Python Binding for MuPDF.” Zenodo. <https://doi.org/10.5281/zenodo.10629631>.
- Melville, James. 2025. *UWOT: The Uniform Manifold Approximation and Projection (UMAP) Method for Dimensionality Reduction*. <https://doi.org/10.32614/CRAN.package.uwot>.
- Meta. 2021. “Harmful Content Can Evolve Quickly. Our New AI System Adapts to Tackle It.” *Meta AI*. <https://ai.meta.com/blog/harmful-content-can-evolve-quickly-our-new-ai-system-adapts-to-tackle-it/>.
- Miller, Marc. 2026. *Safe Social Media Act*. <https://www.canada.ca/en/canadian-heritage/services/safe-social-media-act.html>.
- Ontario. 2026. “Human Development and Sexual Health Education by Grade.” <https://www.ontario.ca/document/health-and-physical-education-grades-1-8/human-development-and-sexual-health-education-grade>.
- OpenAI, :, Aaron Hurst, Adam Lerer, Adam P. Goucher, Adam Perelman, Aditya Ramesh, et al. 2024. “GPT-4o System Card.” <https://arxiv.org/abs/2410.21276>.
- Pariera, Katrina L. 2016. “Barriers and Prompts to Parent-Child Sexual Communication.” *Journal of Family Communication* 16 (3): 277–83. <https://doi.org/10.1080/15267431.2016.1181068>.
- Perring, Christian. 2005. “Review of ‘It’s Perfectly Normal.’” *Northern Wyoming Mental Health Center*. https://wyomentalhealth.org/poc/view_doc.php?type=book&id=2483&cn=28.

- Python Software Foundation. 2026. *Python Language Reference, Version 3.13*. <https://www.python.org>.
- R Core Team. 2025. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Santelli, John, Mary A. Ott, Maureen Lyon, Jennifer Rogers, Daniel Summers, and Rebecca Schleifer. 2006. “Abstinence and Abstinence-Only Education: A Review of u.s. Policies and Programs.” *Journal of Adolescent Health* 38 (1): 72–81. <https://doi.org/https://doi.org/10.1016/j.jadohealth.2005.10.006>.
- SEICUS. 2025. “SEICUS State Profiles.” *SEICUS*. <https://siecus.org/siecus-state-profiles/>.
- Seitz, Amanda. 2022. “Social Media Platforms Fail to Stop LGBTQ Hate Speech, According to GLAAD Report.” *PBS*. <https://www.pbs.org/newshour/nation/social-media-platforms-fail-to-stop-lgbtq-hate-speech-according-to-glaad-report>.
- Singh, Aaditya, Adam Fry, Adam Perelman, Adam Tart, Adi Ganesh, Ahmed El-Kishky, Aidan McLaughlin, et al. 2026. “OpenAI GPT-5 System Card.” <https://arxiv.org/abs/2601.03267>.
- Slowikowski, Kamil. 2024. *Ggrepel: Automatically Position Non-Overlapping Text Labels with 'Ggplot2'*. <https://doi.org/10.32614/CRAN.package.ggrepel>.
- Stryker, Cole. n.d. “What Is Text Classification?” *IBM*. <https://www.ibm.com/think/topics/text-classification>.
- Tavrow P, Guevara F, Lee J. 2025. “How Hatred of Abortion Providers Is Propagated in Social Media: An Investigation of YouTube Videos.” <https://pmc.ncbi.nlm.nih.gov/articles/PMC12573567/>.
- Team, The Pandas Development. 2020. “Pandas-Dev/Pandas: Pandas.” Zenodo. <https://doi.org/10.5281/zenodo.3509134>.
- Tinsley, E. J. 2005. “Immaculate Conception of the Blessed Virgin Mary.” https://books.google.ca/books?id=PN7UMUTBBPAC&pg=PA286&redir_esc=y#v=onepage&q&f=false.
- Tremblay, Chris. 2026. “Transgender Creators Confront 'Pink Shadow Ban' as TikTok and Meta Face Scrutiny over Content Suppression.” *Edge*. <https://www.edgemedianetwork.com/story/168090/transgender-creators-confront-pink-shadow-ban-as-tiktok-and-meta-face-scrutiny>.
- UNESCO, UNFPA, UNAIDS. 2018. “International Technical Guidance on Sexuality Education.” <https://www.who.int/publications/m/item/9789231002595>.
- Venka Simovska, Professor, Dr Ros Kane, Christina R. Peter, Timothy B. Tasker, and Stacey S. Horn. 2015. “Parents’ Attitudes Toward Comprehensive and Inclusive Sexuality Education: Beliefs about Sexual Health Topics and Forms of Curricula.” *Health Education* 115 (1): 71–92. <https://doi.org/10.1108/HE-01-2014-0003>.
- Verena Klein, Rotem Kahalom, Eden Kosman. 2024. “Devaluation of Women’s Sexual Pleasure: Role of Relationship Context and Endorsement of the Madonna-Whore Dichotomy.” <https://doi.org/10.1007/s11199-023-01424-3>.
- Wickham, Hadley. 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D’Agostino McGowan, Romain François, Garrett Grolemund, et al. 2019. “Welcome to the tidyverse.” *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.
- Wickham, Hadley, Joe Cheng, Aaron Jacobs, Garrick Aden-Buie, and Barret Schloerke. 2026.

- Ellmer: Chat with Large Language Models*. <https://ellmer.tidyverse.org>.
- Wickham, Hadley, Jim Hester, and Jennifer Bryan. 2024. *Readr: Read Rectangular Text Data*. <https://doi.org/10.32614/CRAN.package.readr>.
- Woolf, Marie. 2022. “MPs Rush over 150 Amendments to Streaming Bill, Senate Says It Won’t Be Pressured | CBC News.” *CBC*. <https://www.cbc.ca/news/politics/streaming-bill-committee-amendments-1.6490285>.
- World Health Organization. 2026. “Comprehensive Sexuality Education.” *World Health Organization*. <https://www.who.int/news-room/fact-sheets/detail/comprehensive-sexuality-education>.
- Xie, Yihui. 2014. “Knitr: A Comprehensive Tool for Reproducible Research in R.” In *Implementing Reproducible Computational Research*, edited by Victoria Stodden, Friedrich Leisch, and Roger D. Peng. Chapman; Hall/CRC.
- York, Jillian. 2024. “The Global Suppression of Online LGBTQ+ Speech Continues.” *EFF*. <https://www.eff.org/deeplinks/2024/06/global-suppression-online-lgbtq-speech-continues>.
- Zhang, Xinyu, and Wenxue Zou. 2025. “Shattering the Silence: Exploring the Role of Chinese Online Sexual Health Influencers in Promoting Sex Education.” *Sex Education* 25 (1): 95–110. <https://doi.org/10.1080/14681811.2023.2272734>.
- Zhu, Hao. 2024. *kableExtra: Construct Complex Table with 'Kable' and Pipe Syntax*. <https://doi.org/10.32614/CRAN.package.kableExtra>.